

Style-Aware Vision Caption Agent

AMD Developer Hackathon — Track 2: Video Captioning

Four stylistic captions from a two-step vision pipeline

Fireworks AI · MiniMax M3 · Kimi K2P6 · Docker · linux/amd64

Pipeline

1. Download clip & extract 3–6 scene-aware keyframes
2. MiniMax M3 writes a structured visual brief
3. MiniMax M3 verifies the brief against the frames
4. Kimi K2P6 generates four style captions sequentially
5. Keyword guardrails retry any weak style match

Submission Details

Docker Image:

`pjmorales04/amd-track2-agent-v2:latest`

Platform: `linux/amd64`

API: Fireworks AI (MiniMax M3 + Kimi K2P6)

GitHub: `pjmorales1123/amd-track2-agent-v2`